

ANALYSIS:

I tried many kinds of hyperparameters while training the neural network for this situation. The best accuracy i could reach up to was 97.58 percent and this was obtained with a 4 layer(2 hidden) neural network with ReLU activation function and Adam optimizer with a learning rate of 0.0001. All the models i trained were with 30 epochs and 50 batch size.

First, lets see the changes with no of change of layers in the neural network:


The accuracy remained almost same for 1,2 and 3 hidden layers with values of 97.32,97.58 and 97.58 repectively. For no hidden layer too, the network was pretty accurate with 96 percent accuracy, but when u make it 0 layers, i.e basically a regression, the accuracy comes down to around 51 percent. This is because regression requires a lot more training time and isnt very efficient as neural networks.

Next, I tried using different optimizers to see how the result differs, here using SGD gave around 95 percent accuracy and AdaGrad gave around 89 percent, both lesser than Adam which gave a whooping 97.32 percent accuracy.

Now, lets see how the accuracy changed with the change in activation function. ReLU function was more accurate with 97.32 compared to 94 percent accuracy in sigmoid function. This is why ReLU is more widely used than sigmoid nowadays.

Lastly, i tried altering the learning rate and seeing how the result changed. Learning rates of 0.001,0.0001 and 0.00001 gave pretty good accuracies of 96.32,97.32 and 96.0 respectively. As i lessened the learning rate more to 0.000001, accuracy dropped to 89 and when i increased it to 0.01 or more, it dropped drastically to 9.8 percent due to overfitting.

Graph of loss vs no of epochs for sigmoid activation function



Graph of loss vs no of epochs for ReLU activation function,
Note how it decays faster then the one with sigmoid.

